

Bypassing Kernel Barriers

Fuzzing Linux Kernel in Userspace with LKL

Eugene Rodionov, Xuan Xing

6/26/2025



Google

A conversation happened long time ago...

- My manager:
 - Maybe we should fuzz the Linux HID driver?
 - Any idea how to do it?
- <if I were an experienced kernel hacker>:
 - syzkaller?
- Actual me:
 - clang -fsanitize=fuzzer drivers/hid/hid-core.c kernel_stubs.c
- Clang:
 - fatal error: too many errors emitted, stopping now [-ferror-limit=]
- Matt Alexander (former Googler/RedTeamer):
 - Oh, you want to run kernel code in userspace? try LKL!

Agenda

- Introduction
- LKL Basics
- Fuzzing on LKL
- LKL Fuzzers
- Quick Demo
- Conclusion
- Q&A



Introduction

Who are we?

Protect the Android ecosystem through offensive security

- Offensive Security Reviews to verify (break) security assumptions
- Scale through tool development (e.g. continuous fuzzing)
- Develop proof of concepts to demonstrate real-world impact
- Assess the efficacy of security mitigations





But what is LKL?


LKL Overview

- Linux Kernel Library (LKL)¹ builds Linux kernel as a user-space library
 - Implemented as Linux arch-port
 - LKL vs UML
- LKL building blocks
 - Host environment API -- portability layer
 - Linux kernel code
 - LKL syscall API exposed to the user-space application
- Run kernel code without launching a VM
 - Kernel unit testing
 - Fuzzing!^{2,3}

[1] <https://github.com/lkl/linux>

[2] Xu et al., Fuzzing File Systems via Two-Dimensional Input Space Exploration

[3] <https://github.com/atrosinenko/kbdysch>



Using LKL from your C program

```
int ret = lk1_start_kernel(&lk1_host_ops, "mem=50M");

lk1_mount_fs("sysfs");
lk1_mount_fs("proc");
lk1_mount_fs("dev");

int binder_fd = lk1_sys_open("/dev/binder", O_RDWR | O_CLOEXEC, 0);
void *binder_map = lk1_sys_mmap(NULL, BINDER_VM_SIZE,
                                PROT_READ, MAP_PRIVATE, binder_fd, 0);

struct lk1_binder_version version = { 0 };
ret = lk1_sys_ioctl(binder_fd, LKL_BINDER_VERSION, &version);
```

Using LKL from your C program

```
int ret = lkl_start_kernel(&lkl_host_ops, "mem=50M");

lkl_mount_fs("sysfs");
lkl_mount_fs("proc");
lkl_mount_fs("dev");

int binder_fd = lkl_sys_open("/dev/binder", O_RDWR | O_CLOEXEC, 0);
void *binder_map = lkl_sys_mmap(NULL, BINDER_VM_SIZE,
                                PROT_READ, MAP_PRIVATE, binder_fd, 0);

struct lkl_binder_version version = { 0 };
ret = lkl_sys_ioctl(binder_fd, LKL_BINDER_VERSION, &version);
```

Using LKL from your C program

```
int ret = lkl_start_kernel(&lkl_host_ops, "mem=50M");

lkl_mount_fs("sysfs");
lkl_mount_fs("proc");
lkl_mount_fs("dev");

int binder_fd = lkl_sys_open("/dev/binder", O_RDWR | O_CLOEXEC, 0);
void *binder_map = lkl_sys_mmap(NULL, BINDER_VM_SIZE,
                                PROT_READ, MAP_PRIVATE, binder_fd, 0);

struct lkl_binder_version version = { 0 };
ret = lkl_sys_ioctl(binder_fd, LKL_BINDER_VERSION, &version);
```

Using LKL from your C program

```
int ret = lkl_start_kernel(&lkl_host_ops, "mem=50M");

lkl_mount_fs("sysfs");
lkl_mount_fs("proc");
lkl_mount_fs("dev");

int binder_fd = lkl_sys_open("/dev/binder", O_RDWR | O_CLOEXEC, 0);
void *binder_map = lkl_sys_mmap(NULL, BINDER_VM_SIZE,
                                PROT_READ, MAP_PRIVATE, binder_fd, 0);

struct lkl_binder_version version = { 0 };
ret = lkl_sys_ioctl(binder_fd, LKL_BINDER_VERSION, &version);
```

LKL Quirks

- MMU implementation
 - Inspired by MMU implementation in UML
 - Implements 3-level page tables & flat memory model
 - Uses the host's `mmap` to map/unmap LKL's physical pages at provided virtual addresses
- Emulates multiple processes
 - Implements custom `new_thread_group_leader` syscall to modify `tgid`

Applications

- Built-in applications
 - [fs2tar](#): converting file system image to a tar archive
 - [cptofs](#): copies files to/from a file system image
 - [lklfuse](#): mounting filesystem images in userspace, without root privileges
- Kernel fuzzing
 - [kBdysch](#): A collection of kernel fuzzers in userspace
 - [Janus](#): Fuzzing file system using lkl
 - [Sample fuzzers](#): Sample fuzzers developed by Android Red Team (HID, binder)
- Other interesting ideas
 - [LKL.js](#): running Linux kernel on JavaScript directly
 - [User Space TCP](#): a full TCP stack in userspace for better security



Enabling Fuzzing on LKL

Anatomy of LKL fuzzer

LKL enables fuzzing Linux kernel code in user-space


- Use in-process fuzzing engine, such as **libFuzzer**

Advantages

- High fuzzing performance on x86_64
- Ease of custom modifications
 - e.g. mocking hardware, custom scheduler(?)

Limitations

- No SMP in LKL
- x86_64 vs aarch64 -- potential false positives, false negatives



LKL fuzzing limitations

- No SMP support
 - Not easy to test certain concurrency scenarios
- Challenges with coverage build
 - Large binary size
 - Taking too long to build
- Only support x86-64 arch
 - Limited AArch/AArch64 support
 - Drivers specific to non-x86-64 architecture can't be fuzzed
- Need to mock lower level interfaces
 - Mocking hardware support
 - Using virtio backend
- Not in linux main branch (yet)
 - Requiring effort to be upstreamed




LKL Fuzzers

LKL fuzzers


- Human Interface Device (HID) fuzzer
 - A sample LKL fuzzer targeting HID driver
- Binder fuzzer
 - Targeting Android Binder driver
- Virtio driver fuzzing
 - Used by Android pKVM with special attacking scenario (malicious host against guest VM)
 - Virtio based device drivers

Human Interface Device (HID) fuzzing

- Target:
 - Linux HID system
 - Complex spec ([hid1_11.pdf](#), [hut1_5.pdf](#))
- Attacking scenario:
 - Physical access ([juice jacking](#))
- Details
 - Inspired by [kBdysch](#) fuzzing project
 - A proof of concept for LKL fuzzing support development
- Issues discovered and verified with USB simulation
 - Two bulletin class issues (CVE-2020-0465)
 - 6 stability issues



Human Interface Device (HID) fuzzing (cont.)



Exploitability of HID bugs

- We didn't attempt to exploit identified issues in HID driver
- However, ...

The image shows a presentation slide for a Black Hat Europe 2021 briefing. The slide has a yellow header section and a white footer section. The header features the Black Hat logo (a stylized person wearing a hat) and the text "black hat EUROPE 2021" followed by "november 10-11, 2021" and "BRIEFINGS". The footer contains the RISCURE logo ("riscure" in a stylized font) and the tagline "driving your security forward". Below the footer, the title of the briefing is "Achieving Linux Kernel Code Execution Through A Malicious USB Device", presented by "Martijn Bogaard Principal Security Analyst" and "Dana Geist Senior Security Analyst". The URL for the presentation is provided at the bottom.

<https://i.blackhat.com/EU-21/Thursday/EU-21-Bogaard-Geist-Achieving-Linux-Kernel-Code-Execution-Through-A-Malicious-USB-Device.pdf>



DEMO: USB HID Driver issue discovery




Android Binder Driver Fuzzing

What is Binder?

- Primary inter-process communication (IPC) channel on Android
- Support passing file descriptors, objects containing pointers, etc.
- Composed of a userspace library (libbinder) and a kernel driver (/dev/binder)
- Provide Remote Procedure Call (RPC) framework for Java and C/C++


Binder Threat Model



Binder Fuzzing Challenges


Data dependencies

- Binder commands
- Scatter-gather data structures (BINDER_TYPE_PTR)



State dependencies

- Synchronous IPC
 - Cannot send transaction to oneself
 - Multiple outstanding transactions not allowed




Binder Fuzzing Challenges

State dependencies


- Some inputs depend on previous IOCTL calls
 - e.g. Transaction buffers (BC_FREE_BUFFER)

```
ioctl(binder_fd, BINDER_WRITE_READ, x) // 1.  
// y = x->read_buffer->...->buffer  
ioctl(binder_fd, BC_FREE_BUFFER, y) // 2.
```




Multi-process coordination

- All communication requires a **Context Manager**
- Node & Ref setup is required to establish a connection



LKL Binder Fuzzing Harness

- Simulate IPC interactions between multiple clients
 - 3 **clients** (1 Context Manager)
 - IOCTL calls** and **data**



Fuzz data


```
client_1 {
    binder_write {
        binder_commands {
            transaction {
                binder_objects { binder { ptr: 0xbeef } }
            }
        }
    }
}
client_2 {
    binder_read { ... }
    binder_write {
        binder_commands { free_buffer { ... } }
    }
}
client_3 { ... }
client_2 { ... }
client_3 { ... }
client_1 { ... }
```

LKL Randomized Scheduler

Deterministically simulate thread interleaving based on fuzz data¹

Insert yield points before/after synchronization primitives

- spin_lock, spin_unlock
- mutex_lock, mutex_unlock



[1] Williamson, N., Catch Me If You Can: Deterministic Discovery of Race Conditions with Fuzzing. Black Hat USA, (2022).

LKL Binder Fuzzing Results

- Achieved 68% line coverage
- Discovered CVE-2023-20938 & CVE-2023-21255

Coverage Report

Created: 2024-05-06 09:49

Click [here](#) for information about interpreting this report.

| Filename | Function Coverage | Line Coverage | Region Coverage | Branch Coverage |
|-----------------------------------|-------------------------|---------------------------|----------------------------|---------------------------|
| binder.c | 79.14% (110/139) | 68.55% (3086/4502) | 50.90% (4486/8813) | 50.50% (1317/2608) |
| binder_alloc.c | 78.38% (29/37) | 70.41% (564/801) | 49.87% (752/1508) | 49.20% (185/376) |
| binder_alloc.h | 50.00% (1/2) | 11.11% (1/9) | 50.00% (1/2) | - (0/0) |
| binder_internal.h | 60.00% (3/5) | 68.75% (11/16) | 73.68% (14/19) | - (0/0) |
| Totals | 78.14% (143/183) | 68.73% (3662/5328) | 50.79% (5253/10342) | 50.34% (1502/2984) |

Generated by llvm-cov -- llvm version 12.0.6git




DEMO: Test case for CVE-2023-20938




Virtio fuzzing for Android pKVM

Android Protected KVM Attack Surface



Protected VM virtio attack surface



Virtio Front-end Fuzzers

Kernel under test:

- Android13-5.10

virtio_ring:


- fuzzes ring-buffer processing functionality
- handles both split & packed mode

virtio_pci:

- fuzzes PCI configuration space
- LKL arch-specific implementation of PCI bus
- mock-out PCI MMIO in the fuzzer harness

virtio_blk:

- mutates the virtio_blk configuration block



Virtio_blk fuzzer finding

```
int block_read_full_page(struct page *page, get_block_t *get_block)
{
    struct buffer_head *bh, *head, *arr[MAX_BUF_PER_PAGE];

    ...
    do {
        if (buffer_uptodate(bh))
            continue;

        if (!buffer_mapped(bh)) {
            int err = 0;

            ...
        }
        arr[nr++] = bh;      ===== OOB write on stack
    } while (i++, iblock++, (bh = bh->b_this_page) != head);
    ...
}
```

Virtio_blk fuzzer finding

- With the block size **0xe5e5e5e5**:
 - `inode->i_blkbits == 32`
 - `1 << READ_ONCE(inode->i_blkbits)` is undefined behavior in C
 - `1 << READ_ONCE(inode->i_blkbits) == 1` on x86 architecture

```
static struct buffer_head *create_page_buffers(struct page *page, ...)  
{  
    BUG_ON(!PageLocked(page));  
    if (!page_has_buffers(page))  
        create_empty_buffers(page, 1 << READ_ONCE(inode->i_blkbits), b_state);  
    return page_buffers(page);  
}
```



Conclusions

Key takeaways

- LKL can be an effective way to fuzz certain kernel scenarios
 - Scenarios requiring complex setup (e.g. hardware-to-kernel attack surface)
 - Dependencies to some userspace libraries
 - Visual code coverage and debugging can be very useful
 - It's good to have alternatives
- The fact that you don't know something might be your advantage
 - Discovering new approaches from naive ideas
 - Learning without predefined mindset

Future work

- Enable MMU (<https://github.com/lkl/linux/pull/551>)
- Upstream Android Binder fuzzer (<https://github.com/lkl/linux/pull/564>)

Creating more fuzzers and profit!

Acknowledgement

- Octavian Purdila and LKL maintainers
- <https://github.com/atrosinenko/kbdysch>
- Android pKVM team
- Android Binder team

Resources

- [libFuzzer – a library for coverage-guided fuzz testing](#)
- [Achieving Kernel Code Execution through Malicious USB device](#)
- [Android Virtualization Framework and pKVM security](#)
- [Emulating USB device with Python](#)
- [Attacking Android Binder: Analysis and Exploitation of CVE-2023-20938](#)
- [Binder Internals](#)



Questions?

